

PlanMyAgents — Business Plan

The trust, recipe, and marketplace platform for the AI agent economy.

Version: 2.1 (May 2026 — post-pivot to discovery + recipe + marketplace + partnerships) **Stage:** Pre-revenue, pre-incorporation, working prototype with 1,200+ backend tests passing **Ask:** \$3–5M seed, 18-month runway, Series A readiness in 12 months

Table of contents

1. [Executive summary](#)
 2. [The market shift](#)
 3. [Customer and pain — both sides](#)
 4. [Product — the three-phase ladder](#)
 5. [Strategic positioning](#)
 6. [Business model — the 9-line revenue ladder](#)
 7. [Marketplace economics \(Phase 2\)](#)
 8. [Partnership ecosystem \(Phase 3\)](#)
 9. [Go-to-market](#)
 10. [Competitive landscape](#)
 11. [Moat and defensibility](#)
 12. [Product roadmap \(24 months\)](#)
 13. [Financial plan](#)
 14. [Team and hiring](#)
 15. [Risks and mitigations](#)
 16. [Fundraising narrative](#)
 17. [Appendix: glossary and FAQ](#)
-

1. Executive summary

The market shift. Between late 2024 and mid 2026, three things happened: Anthropic shipped MCP, Google et al shipped A2A, and every major LLM client (Claude Desktop, Cursor, Cline, ChatGPT, Gemini) added native tool use. The result is a fragmented, exploding supply of AI agents — over 1,000 MCP servers, hundreds of A2A agents, thousands of AI-native services per month being shipped. Nobody ranks them. Nobody verifies claims. Nobody runs a vendor-neutral marketplace where vendors can be discovered and buyers can find trusted agents.

The opportunity. Every dense supply ecosystem produces a discovery + marketplace layer worth billions: Hugging Face for models (~\$4.5B), G2 for SaaS (~\$2B+), Capterra (now part of \$25B+ Gartner Digital Markets), npm for packages, App Store for apps, Steam for games. The agent ecosystem is 10x the surface area of any of those — and its discovery + marketplace + partnership layer is empty.

Our product is a three-phase platform:

- 1. **Phase 1 — Plan + Discover + Recipe** (*now* → *Y1*): Decompose goals into sub-tasks, discover candidate providers, rank them with evidence, export recipes when steps are actually exportable, and refuse with reasons when gates fail.
- 2. **Phase 2 — Marketplace** (*Y1 end* → *Y2*): Vendors claim listings, get tested benchmarks, run sponsored category placements (with full disclosure), pay for demand-data API access.
- 3. **Phase 3 — Partnership Platform** (*Y2* → *Y4*): OEM into n8n / Zapier / Claude Desktop / Cursor / Cline; private enterprise registries; standards influence; cloud platform channels.

The proprietary engine under all three phases is **Planning + Discovery** — LLM goal decomposition, embedding-driven candidate matching, trust ladder verification, LLM judging, ranker. **Recipe + Benchmark + Marketplace** are how we monetize the engine.

The wedge is developer agent workflows — densest MCP/A2A supply, fastest validation, viral via dev communities — followed by GTM data workflows and then horizontal expansion.

The moat is six compounding assets: discovery freshness × verification evidence × benchmark history × refusal demand data × **marketplace network effects** × **Tier-1 partnership distribution**.

The ask. \$3–5M seed funds 18 months of: shipping recipe export, reaching 25K weekly active developers, signing **Tier-1 partnerships with n8n + Cursor + Anthropic** (distribution + dev mindshare + protocol authority), running a parallel **founding marketplace cohort workstream** to land 5–10 paying vendors with first 3 publishable Verified Benchmarks, closing 3 enterprise design partners, and reaching \$50K MRR — the Series A gate.

Anti-pitch. We are *not* Zapier for agents — we partner with Zapier and n8n, we don't compete on connector count. We are *not* an executor with stored credentials — we never hold customer keys. We are *not* a pay-to-rank marketplace — vendors pay for visibility, never for ranking position. We do, in Phase 2, offer an **optional 1–2% vendor-opt-in execution fee** (Amazon Associates–style attribution, paid by the vendor out of their margin, user price never changes, capped at 2%, decoupled from ranking) — this is a vendor-side attribution mechanism, not a universal take rate. The first two anti-positions remove the connector-arms-race trap; the third (with the execution-fee cap) is the trust firewall that makes us defensible as we scale.

2. The market shift

2.1 The 14-month supply explosion

Timeline	Event	Supply impact
Nov 2024	Anthropic launches MCP	First open standard for LLM tool integration
Q1 2025	First MCP registries (Smithery, MCP Marketplace) appear	~200 servers
Q2 2025	A2A protocol announced (Google, Salesforce, SAP, ServiceNow)	First agent-to-agent interop standard

Timeline	Event	Supply impact
Q3 2025	OpenAI Apps + GPT function-calling refresh	Multi-vendor tool standardization
Q4 2025	Cursor, Cline, Continue.dev all ship native MCP	Buyer surface explodes
Q1 2026	Official MCP Registry launches	Quality bar raised, supply still fragmented
Q2 2026	n8n + Make + Zapier add AI / MCP nodes	Workflow tools acknowledge they need an upstream discovery layer they can't build
Now	1,000+ MCP servers; 200+ A2A cards; 1,000+ AI-native services	No trusted ranking, no vendor-neutral marketplace, no partnership API

2.2 Why the discovery + marketplace layer never emerged organically

Every comparable supply ecosystem produced a discovery + marketplace layer within 3–5 years of standardization:

Ecosystem	Standardization	Discovery + marketplace layer
Open-source packages	npm 2010	npm registry (2010), libraries.io
ML models	TensorFlow 2015 / PyTorch 2016	Hugging Face Hub (2018)
SaaS	REST + OAuth ~2010	G2 (2012), Capterra (2010s), Gartner Digital Markets
Mobile apps	iOS 2008	App Store launch-day, Play Store
Steam games	Steamworks ~2003	Steam Store with reviews & rankings
Browser extensions	Chrome Web Store 2011	Chrome Web Store

The MCP/A2A ecosystem is at the equivalent of "npm in 2011 plus G2 in 2012" — supply is dense, supplier behavior is unaligned, and **no trusted ranking + vendor-neutral marketplace exists**. The first credible layer captures category-defining brand value.

2.3 What we are betting on (and what would invalidate)

We are betting on:

- Continued MCP / A2A adoption (already exponential)
- Multi-model, multi-LLM-host reality (no single vendor wins all)
- Buyer pain around discovery being severe enough to pay for trust
- **Vendors paying to be discovered and benchmarked** (G2 / Capterra precedent)
- **Workflow tools and LLM hosts wanting upstream discovery they don't have to build** (n8n, Cursor, Claude Desktop precedent)

What would invalidate:

- A frontier model vendor launches a *vendor-neutral, multi-host, multi-protocol* marketplace (defense: very unlikely; their incentive is single-host capture)
 - MCP gets replaced by a closed-vendor standard (defense: we already cover A2A and AI-native services in parallel)
 - The agent ecosystem stalls — agents fail to deliver real workflow value (this is the real existential risk; pace of progress in early 2026 suggests unlikely)
-

3. Customer and pain — both sides

The product serves a two-sided market: **users** (Phase 1) and **vendors** (Phase 2 onward). Pain on each side is severe and distinct.

3.1 Demand-side ICP — three tiers, sequenced

Tier 1 — Primary (Year 1): Developer / agent builder

- Builds with Claude Desktop, Cursor, Cline, Continue.dev, or own LLM application
- Wires 5–30 MCP servers / A2A agents into their workflow
- 2–4 hour weekly cost of hand-evaluating new agents
- Pays for dev tools (\$20–100 / month willingly)
- Viral via HN, dev Twitter, Reddit, agent-dev Discords

Tier 2 — Secondary (Year 2): Knowledge worker / operator (RevOps, recruiter, analyst)

- Already uses Clay, Apollo, Hunter, Bardeen, Apify
- Wants guided agent workflows without writing code
- Pays company card (\$50–200 / user / month)
- Cares deeply about output quality and verifiable cost

Tier 3 — Tertiary (Year 2–3): Enterprise tech / platform team

- Asked to "standardize on which agents we approve internally"
- Needs private registry, governance, on-prem, SLA
- Procurement cycle \$50K–500K

3.2 Supply-side ICP — three vendor archetypes

Archetype A — Independent agent vendor (Apollo, Hunter, Perplexity, Linkup, Tavily, Apify, Browserbase, Exa)

- Wants vendor-neutral validation to win enterprise deals
- Wants exposure across all LLM hosts and workflow tools at once
- Today: lists self in 7 catalogs; no neutral benchmark; can't measure demand
- Will pay: claimed profile, verified benchmark badge, sponsored placement, demand-data API

Archetype B — Open-source MCP author (community devs publishing MCP servers)

- Wants their MCP discovered + adopted

- Doesn't have budget — wants free tier and earns badge through quality
- Will benefit from: free claimed profile, free probe-based capability_verified status
- Their adoption fuels the marketplace's "free tier credibility" without revenue immediately

Archetype C — Enterprise vendor with public API (Snowflake, Databricks, Datadog, Stripe, Zendesk)

- Will eventually publish official MCP servers
- Wants their MCP to be *trusted* the moment it ships
- Will pay: enterprise verified benchmark, sponsored placement, lead routing, integration into our enterprise registry

3.3 Demand-side buyer journey today (what we replace)

Status quo workflow for a developer adding a new capability to their LLM app:

1. Google "best MCP for X" → 5 conflicting blog posts, 2 outdated
2. Browse Smithery / MCP Marketplace → 30 results, no ranking, no benchmark
3. Pick one, install, test in Claude Desktop → 30% chance it works first try
4. If it breaks, repeat with fallback → another hour
5. Compare costs / latency manually → never done because tedious
6. Write internal Notion doc; six months later it's stale

Total cost per new capability: 2–4 engineer hours plus opportunity cost.

Replacing this: 1 search → ranked top 3 with verification + benchmark → recipe drops straight into config.
Total time: under 5 minutes.

3.4 Supply-side vendor journey today (what we replace)

Status quo for an agent vendor wanting to be adopted:

1. Submit to Smithery, MCP Marketplace, official MCP Registry, GitHub topic, AI directories — 7 separate forms
2. Hope LLM hosts curate yours into Claude Desktop / Cursor docs (they don't, no methodology)
3. Hope someone benchmarks you (no one does at commercial workflow level)
4. Get inbound contact from one buyer; have no demand data to size the market
5. Rebuild your own analytics from referrer logs; never accurate
6. Spend marketing budget on vendor lists and dev community posts; CPL is brutal

Replacing this: one claimed profile → verified benchmark badge → sponsored placement when desired → demand-data API to size markets. Single vendor-neutral marketplace recognized across all LLM hosts and workflow tools.

3.5 Quantified willingness-to-pay

Cohort	Anchor	Willingness to pay
Developer	Time wasted (2–4 hr/wk) + dev tool comp	\$20–50 / user / mo
Operator	Time wasted + Clay/Apollo comp	\$50–200 / user / mo

Cohort	Anchor	Willingness to pay
Enterprise	Gartner Magic Quadrant license comp	\$50K–500K / yr
Vendor (claimed profile)	G2 vendor profile baseline	\$500–2K / yr
Vendor (verified benchmark)	Underwriters Laboratories tested-badge comp	\$5–50K / cycle
Vendor (sponsored placement)	Capterra sponsored listing comp	\$10–100K / qtr
Vendor (demand-data API)	Gartner Research subscription comp	\$10–50K / yr

4. Product — the three-phase ladder

4.1 Phase 1 — Plan + Discover + Recipe (now → Y1)

The proprietary engine + free user tier + Pro tier. Four user-facing surfaces:

Surface	User	What they get	Pricing
/discovery/search	Developer	Ranked candidate list with trust + benchmark badges	Free
/goal → recipe	Developer / operator	Decomposed sub-tasks; best agent per task; recipe download (Claude / n8n / Cursor / markdown)	Free → Pro
/benchmark/{cap}	Vendor-neutral buyer	Programmatic comparison data	Paid API
/open-mcp-opportunities	Builders	Capabilities where APIs exist but no MCP/A2A wraps them	Free

Pro tier (\$20–50 / user / month): saved recipes, team workspaces, freshness alerts, recipe versioning, weekly index digest.

4.2 Phase 2 — Marketplace (Y1 end → Y2)

Vendor-side surfaces. Five revenue lines, one shared firewall (rankings stay neutral; money buys *visibility*, never *ranking position*).

Surface	Vendor offer	Pricing
Claimed agent profile	Claim listing, edit description, add screenshots/docs, respond to community	\$500–2K / yr
Verified Benchmark badge	Extended benchmark (≥250 samples, p95 latency, success rate, cost) — result published whatever it is	\$5–50K / cycle
Sponsored category placement	"Sponsored" badge in category pages; never above natural #1	\$10–100K / qtr

Surface	Vendor offer	Pricing
Lead routing (opt-in)	When user goal needs vendor X's capability, vendor optionally notified (with user consent)	\$10–100 / qualified lead
Demand-Data API	Vendor pulls demand signal data ("what are users asking for that no agent fulfills?")	\$10–50K / yr

4.3 Phase 3 — Partnership Platform (Y2 → Y4)

OEM + enterprise + standards. Distribution becomes the moat.

Surface	Partner	Deal shape
Recipe export integration	n8n, Zapier, Make	Native template export from <i>/recipe</i> ; co-marketing
"Add to host" buttons	Cursor, Cline, Continue.dev, Claude Desktop	One-click recipe install in target host
Joint methodology	Anthropic (MCP), Google (A2A)	Co-published quality methodology; "Tested by PlanMyAgents" badge in host UIs
Cloud OEM	AWS Bedrock, Azure AI Foundry, GCP Vertex	OEM index into their dev experience
Enterprise registry	Fortune 500 platform teams	Private agent registry, on-prem, governance, SLA
Standards influence	MCP WG, A2A WG, OpenAPI Initiative	Implementer-of-record for "trusted discovery"

4.4 Optional execution — BYO-credentials sandbox (across all phases)

A user can choose to execute a recipe step from within the PlanMyAgents UI rather than copying it to their LLM host. **Constraints:**

- Credentials live in user's browser session only (never persisted server-side)
- Execution goes through *GenericProtocolAdapter* for MCP / A2A / OpenAPI only
- No hand-coded vendor wrappers
- Cost cap enforced per session via *spend_ledger.py*
- Every call audited; no telemetry data flows back to vendors
- Gated by env var (off by default in production until trust review)

Why this matters: lets us collect benchmark data on real workflows without becoming an executor. The user is always in the driver seat.

5. Strategic positioning

5.1 What we are

"The trust, recipe, and marketplace platform for the AI agent economy — G2 + Hugging Face + Gartner Digital Markets, for an ecosystem 10× the surface area of any of them."

The vendor-neutral, multi-protocol, multi-LLM-host discovery + ranking + recipe + marketplace + partnership platform for AI agents, MCP servers, and A2A agents.

5.2 What we are deliberately not

Anti-position	Why we refuse
Zapier / n8n / Make competitor	Connector arms race; they're our Tier-1 distribution partners
Workflow builder UI	Layer above us; we feed those tools
Pay-to-rank marketplace	Trust firewall; rankings stay neutral; money buys visibility only
Hand-coded vendor wrapper farm	We index public agents through open protocols only
Executor with stored credentials	Compliance burden, trust risk, no moat
Frontier model competitor	Different problem; we ride on every LLM host
Marketplace with transaction fees on every recipe execution	Compromises BYO-credentials trust; deferred to "vendor-opt-in only" Phase 3 option

5.3 The category we are creating

There is no existing category for "vendor-neutral, benchmarked, recipe-first, two-sided agent marketplace." We are creating it. The mental model we want adopted:

"Before I add a new agent to my workflow, I check PlanMyAgents — like I check npm for packages, Hugging Face for models, G2 for SaaS, App Store for apps."

5.4 Why "trust" is the keyword (vs "discovery" or "marketplace" alone)

Discovery is a feature; Smithery has it. Marketplace is a model; G2 has it. **Trust is the asset that compounds across both sides** — buyers trust us to recommend; vendors trust us to be neutral; LLM hosts trust us as the methodology authority. Without trust, we are one of many catalogs. With trust, we are *the* layer.

6. Business model — the 9-line revenue ladder

Each tier unlocks **after** the previous one has proof. We do not ship all nine on Day 1.

#	Tier	Product	Unlocks after	ARPU / ACV
0	Free	Discovery search, basic recipe, public leaderboards	Day 1	\$0 (acquisition + brand)

#	Tier	Product	Unlocks after	ARPU / ACV
1	Pro	Saved recipes, teams, private notes, recipe versioning, freshness alerts, watchlists	10K weekly active developers	\$20–50 / user / mo
2	Benchmark API	Programmatic ranking + benchmark data	3+ publishable benchmark cells	\$500–5K / mo
3	Enterprise registry	Private registry, on-prem, governance, SLA, audit, SSO	3+ design partners	\$50K–500K / yr
4	Claimed Vendor Profile	Vendor-edit access; analytics; community response	50+ verified agentic candidates per category	\$500–2K / yr per vendor
5	Verified Benchmark Badge	Extended ≥250-sample run; published result whatever it is	1+ publishable benchmark cell live	\$5–50K / cycle per vendor
6	Sponsored Category Placement	Category page sponsored slot (disclosed)	25K+ WAU	\$10–100K / qtr per vendor
7	Demand-Data API	Vendor pulls "what users are asking for that no agent fulfills"	100K+ goal queries/mo	\$10–50K / yr per vendor
8	Optional Execution Fee (<i>vendor-opt-in</i>)	1–2% on calls flowing through our BYO-credentials sandbox to a vendor that has opted in. Vendor pays out of their own margin; user price never changes; capped at 2%; decoupled from ranking; vendor benefits = execution attribution analytics + sponsored eligibility + co-marketing rights	BYO-creds sandbox GA + 5+ paying vendors on tiers 4–6	1–2% per sandboxed call (vendor-borne); typical vendor MRR \$1–25K

6.1 What we deliberately refuse to monetize

- ✗ User-side margin on agent execution (user never sees a price change from anything we do — even tier-8 is vendor-paid)
- ✗ Float on agent execution (we never hold the money flow; vendor settles fee monthly out-of-band)
- ✗ Hidden vendor placement fees (destroys trust)
- ✗ Per-agent listing fees (destroys index density and impartiality)
- ✗ Referral fees that bias ranking (destroys benchmark credibility)

- **✗** Customer data sold to vendors (destroys privacy claim)
- **✗** Pay-to-rank — the vendor-neutrality firewall. Even tier-8 execution fee is **hard-capped at 2% and architecturally decoupled from the ranker**; vendor cannot pay a higher rate to influence position.

6.2 Why the optional execution fee (tier 8) is *not* a "Zapier-style take rate"

Dimension	Zapier-style take rate (rejected)	Our optional execution fee (tier 8)
Who pays	User (margin on top of vendor price)	Vendor (out of their own margin)
User price impact	Higher	Unchanged
Universality	Mandatory on every execution	Opt-in per vendor; default off
Rate	Variable, can rise to 20–30%	Hard-capped at 2% by contract
Ranking impact	Often biased toward higher-fee vendors	Decoupled from ranking by architecture (P11)
Vendor benefit	None (pure cost)	Attribution analytics + sponsored eligibility + co-marketing
Disclosure	Often hidden	Always disclosed in recipe metadata
Comparable	Stripe / Zapier / App Store	Amazon Associates attribution fee

The vendor-side, opt-in, capped, decoupled, vendor-paid framing is what keeps this consistent with the trust position. Vendors opt in because the attribution and co-marketing value is worth the 1–2% to them.

6.3 Indicative unit economics (Year 2 target)

Pro tier:

- ARPU: \$30 / user / month; gross margin: 85%; payback: ~3 months

Benchmark API:

- ARPU: \$1,500 / month average; gross margin: 70%; payback: ~6 months

Marketplace (vendor side, blended):

- ARPU: \$5,000 / vendor / year (mix of claimed profile + verified badge + sponsored)
- Gross margin: 80% (mostly platform compute + benchmark cost)
- Payback: ~9 months on \$4K acquisition cost
- LTV: \$25K+ over 5 years (sticky vendor profiles)

Enterprise registry:

- ACV: \$150K average; gross margin: 75%; sales cycle: 6–9 months

Optional execution fee (tier 8):

- Vendor MRR: \$1–25K per opted-in vendor (depends on sandboxed execution volume to that vendor)
- Gross margin: 95% (no incremental compute; vendor settles monthly)
- Payback: instant (no acquisition cost beyond the existing claimed-profile relationship)
- Adoption assumption: 10% of paying vendors opt in by Year 2 end; 25% by Year 3 end

6.4 Why this is fundable at venture scale

The marketplace + partnership lines are *additive* to the dev-tool subscription lines. Total ARR ramp resembles G2 (2012–2018: \$0 → \$50M) overlaid on Hugging Face (2019–2024: \$0 → \$70M+) — both built on dense, fragmented supply ecosystems. Our supply ecosystem is bigger than either.

Year	ARR target	Headcount	Revenue mix
Y1	\$500K	7	55% Pro, 30% Benchmark API, 10% early marketplace (tiers 4-5), 5% Enterprise
Y2	\$5M	14	30% Pro, 25% Benchmark API, 20% Marketplace (tiers 4-6), 15% Enterprise, 5% Execution fee (tier 8), 5% Partnership royalties
Y3	\$20M	28	22% Pro, 18% Benchmark API, 22% Marketplace, 22% Enterprise, 10% Execution fee, 6% Partnership royalties
Y4	\$50M	55	18% Pro, 17% Benchmark API, 25% Marketplace, 22% Enterprise, 12% Execution fee, 6% Partnership royalties
Y5	\$100M+	110+	Same mix, larger contracts, international

7. Marketplace economics (Phase 2)

Full design in [docs/marketplace-design.md](#). Summary here.

7.1 The vendor-neutrality firewall


This is the asset that distinguishes us from every other "agent registry" attempt. **Money buys visibility (and opt-in execution attribution). Money never buys ranking position.** Concretely:

- **Rankings** are computed only from: cosine relevance × verification tier weight × benchmark score × freshness × cost penalty. **No vendor payment input — including the optional execution fee.**
- **Sponsored placement** lives in a clearly-labeled "Sponsored" section, always *below* the natural #1.
- **Verified Benchmark** results are published whatever they are. A vendor cannot pay to have a poor result hidden — they can only pay for the run; the result is the result.
- **Claimed profile** lets vendors edit *their description and metadata*; never their ranking score.
- **Optional execution fee (tier 8)** is hard-capped at 2% in contract and decoupled from the ranker by architecture — vendor cannot pay a higher rate to influence position.
- All sponsored / paid / opt-in-execution-fee relationships disclosed in-surface and machine-readable in the API at [/disclosure](#).

7.2 Why vendors will pay despite the firewall

Reason	Comparable
Single neutral marketplace recognized across all LLM hosts	G2 cross-category buyer reach
Verified benchmark badge unlocks enterprise procurement	UL / SOC2 / ISO badges
Sponsored category placement reaches high-intent buyers at moment of decision	Capterra sponsored listings
Demand-data API tells vendor what to build next	Gartner Research subscription
Reach without engineering investment in 7 separate catalog submissions	Marketplace consolidation value

7.3 Vendor portal flow (Phase 2)



7.4 The optional vendor-opt-in execution fee (tier 8)

Phase 2 introduces a deliberately **optional, vendor-side, capped, decoupled** execution fee. Most "marketplace take rate" implementations destroy trust because they are user-facing margin, mandatory, opaque, and ranking-influencing. Ours is none of those things by construction:

Invariant	What it means
Opt-in per vendor	Vendor signs an explicit addendum to claim this fee. Default state for every vendor is OFF.
Vendor-paid, never user-paid	The vendor pays us 1–2% out of their own margin. The user pays the vendor's posted price, unchanged. We never sit in the user's payment flow.
Hard rate cap (2%)	Contractually capped. Vendor cannot "pay more for better placement" — the lever doesn't exist.

Invariant	What it means
Decoupled from ranker by architecture	The ranker (<code>agents/router.py</code> + <code>workflows/scoring.py</code>) does not import the marketplace store. Same firewall protecting sponsored placement (P11) protects the execution fee.
Disclosed in recipe metadata	Every recipe step that touches an opted-in vendor carries <code>disclosure.execution_fee_active: true</code> . Listed in <code>/disclosure</code> page.
Triggered only via our BYO-creds sandbox	We must be in the runtime path to claim it. User executing recipe in their own LLM host (Claude Desktop, Cursor, n8n) without our sandbox → no fee.
Vendor benefits = attribution + sponsored eligibility + co-marketing	Vendor opts in because they want the attribution data and co-marketing rights — not because we forced them.

The closest comparable is **Amazon Associates / affiliate attribution**, not Stripe / Zapier / App Store. The check on whether we ever cross into a Zapier-style universal take rate is the 2% cap + opt-in-default-off + vendor-paid trifecta — those three together make the trust position non-degradable.

8. Partnership ecosystem (Phase 3)

Full design in [docs/partnership-strategy.md](#). Summary here.

8.1 Tier-1 — founding-priority partnerships (months 0–12)

Three partnerships covering all three distribution arms. Each one's cost of inaction is highest at this stage.

Partner	Pain we solve	Our gain	Deal shape	Risk
n8n	Their users need agent recommendations they can't produce; native MCP integration emerging	100K+ engaged workflow-builder distribution; legitimacy with workflow-tool buyers	Co-built "Recipes for n8n"; native template export; revenue share on Pro upsells from n8n users	Low (orthogonal product)
Cursor	Their users hand-search MCPs and tool integrations; Cursor docs can't be the trust authority across vendors	Recipe distribution into the largest native-MCP IDE; dev mindshare; viral via Cursor's content	"Add to Cursor" buttons on every recipe; co-published agent recommendation methodology; Cursor docs link to our <code>/search?host=cursor</code> ; eventual <code>cursor.directory</code> integration	Low-medium (their own MCP curation might extend; orthogonal because we are vendor-neutral cross-host)

Partner	Pain we solve	Our gain	Deal shape	Risk
Anthropic	MCP ecosystem is vendor-noisy; they need trust signals; can't be vendor-neutral themselves	Brand association; protocol authority; Claude Desktop distribution	Joint MCP quality methodology; "Tested by PlanMyAgents" badge in Claude Desktop; Anthropic-funded benchmark cells (with disclosure)	Medium (could launch competing registry — co-opt via methodology partnership)

The three together cover all three distribution arms — workflow tools (n8n) + IDE/coding host (Cursor) + protocol authority (Anthropic) — and lock in our position before any of them in-sources the layer.

8.1.1 Founding marketplace cohort — parallel workstream (not a partnership tier)

Distribution partnerships put the engine in front of users. The founding marketplace cohort puts **first paying vendors** on the engine — proving the Phase-2 marketplace economics before public launch.

Aspect	Detail
Target	10 vendors over 12 months (Apollo, Perplexity, Hunter, Linkup, Tavily, Exa, Apify, Firecrawl, Browserbase, Resend)
Selection criteria	Already shipping production MCP/A2A; benchmark-forward; will commit to publishing results; sells to dev/operator ICP
Offer	Founding price 50% off Verified Benchmark Y1; 3-year price lock on all SKUs; product council seat; case study rights; co-marketing on launch; founding cohort preferred adopters of optional execution fee (tier 8) with 1-year fee waiver
What they validate	(a) Vendor-side willingness to pay (\$25K–250K marketplace ARR Y1), (b) First 3 publishable benchmark cells, (c) Optional execution fee viability (3+ opt-ins = tier 8 works), (d) Distribution into their existing communities
Why this is workstream, not partnership tier	Different motion (sales not BD); different metric (vendor ARR not user reach); different cadence (rolling not lumpy); founder + Partnership BD lead jointly own

8.2 Tier-2 — distribution + scale (Year 1–2)

Partner	Strategic shape	Risk
Cursor / Cline / Continue.dev	"Add to host" buttons; doc links to our /search ; co-marketing	Low
Zapier	Co-marketing; recipe export Zapier format; Zapier links from agent help docs	Medium-high competitive — manage by being the <i>neutral</i> layer above their internal recommendations

Partner	Strategic shape	Risk
Google (A2A)	Joint A2A benchmark methodology; A2A-funded benchmarks	Low
OpenAI	Recipe export to OpenAI Apps; benchmark for OpenAI tool quality	Medium (their own marketplace ambitions)
Vercel / Supabase / Replicate / Modal	Recipe deploys as their function; co-marketing	Low
LangChain / LlamaIndex / CrewAI / AutoGen / smolagents	Recipe export targets their framework; they recommend us for discovery	Low
GitHub / Microsoft	OEM into GitHub agent suggestions	Medium-high (natural acquirer)


8.3 Tier-3 — long-arm strategic + standards (Year 2+)

Partner	Why
Cloud platforms (AWS Bedrock, Azure AI Foundry, GCP Vertex)	OEM into developer experience; enterprise channel partner
Gartner / Forrester / IDC	Industry analyst legitimacy; would pay us for data
Linux Foundation AI / CNCF	Open neutrality positioning
MCP WG / A2A WG / OpenAPI Initiative	Standards influence
Auth-for-agents (Clerk / Auth0 / WorkOS / Styth)	Recipe needs OAuth flows for vendor APIs
Observability (Datadog / Honeycomb / Helicone / Langfuse)	Telemetry handoff; "which agent broke"
Universities (Stanford, MIT, Berkeley AI labs)	Academic credibility for benchmark methodology; talent pipeline

8.4 Partnerships we deliberately decline

Suggested	Why we decline
Stripe / Plaid / payment vendors as adapter customers	Conflicts with "no execution" philosophy; would push us toward executor model
Salesforce / HubSpot direct integrations at seed	Too far from agent-builder ICP; comes naturally via enterprise
Vector DB vendors (Pinecone, Weaviate) as integration partners	Different layer, no buyer overlap

8.5 The partnership flywheel



9. Go-to-market

9.1 Sequenced beachhead

- 1. **Months 0–6: Developer agent workflows.** Acquisition through HN, dev Twitter, Reddit, GitHub presence, agent-dev Discords. Pricing: free with optional Pro. Distribution gain: **n8n partnership signed by month 6** for recipe export integration.
- 2. **Months 6–12: Anthropic methodology partnership announced.** First publishable benchmark cell live. First 5 vendor customers paying for claimed profile. GTM data category content expansion.
- 3. **Year 2: Vendor portal GA + Cursor / Cline / Continue partnerships.** Enterprise registry pilots with 3 design partners.
- 4. **Year 2–3: Cloud OEM (AWS / Azure / GCP).** Sponsored placement program scales. Demand-data API enterprise customers.
- 5. **Year 3+: Horizontal expansion** to any category with measurable workflow outputs.

9.2 Acquisition channels by stage

Stage	Primary channel	Secondary	Cost model
-------	-----------------	-----------	------------

Stage	Primary channel	Secondary	Cost model
Developer	HN, dev Twitter, Reddit, GitHub repo, agent-dev Discords	Tech publications, MCP/A2A community events	Content + DevRel
Operator	SEO ("best agents for X"), Clay/Apollo adjacent content	Partner referrals, RevOps newsletters	Content + paid SEO
Vendor	Direct outreach to vendor product / RevOps teams	Inbound from vendor-neutral content; analyst recognition	BD + content
Enterprise	Inbound from Pro users with corporate cards	Conference (KubeCon, AI Engineer Summit, Gartner shows)	Sales
Workflow tool / LLM host	Direct partnership BD	Conference + product-integration co-marketing	BD

9.3 Founding distribution motion (months 0–6)

- **Public launch:** HN front page; product-first post ("Show HN: G2 + Hugging Face for the agent economy")
- **GitHub presence:** OSS components (curated registries, embedding utilities, scout SDK template) for credibility
- **Content cadence:** 1 deep "best agents for X" post / week; doubles as SEO and lead magnet
- **Community engagement:** Active in MCP Discord, A2A Slack, agent-dev Twitter; ship recipes for community members publicly
- **Partner BD outreach:** n8n, Cursor, Cline, Continue, Anthropic — first conversations in month 1
- **Founding vendor cohort:** Apollo, Perplexity, Hunter, Linkup, Apify — invited to "founding partner" program with extended free claimed profile + discounted Verified Benchmark
- **Design partner program:** 5 dev teams + 3 enterprise early access cohorts with direct founder support

10. Competitive landscape

10.1 Direct competitors (today)

Player	What they do	Their moat	Our edge
Smithery	MCP catalog with search	Index volume, head start	We rank, recipe, benchmark, run a vendor-neutral marketplace; we are <i>one layer above</i> (we use them as input)
MCP Marketplace	MCP catalog	Vendor relationships	Same as above; protocol-neutral (also A2A + AI-native)
Official MCP Registry	Protocol-official catalog	Anthropic association	We rank and benchmark; they only list. Anthropic is our Tier-1 partner.

Player	What they do	Their moat	Our edge
AgentBench / AgentSearchBench / Steel.dev	Research benchmarks	Academic rigor	We focus on commercial workflow execution + marketplace

10.2 Adjacent / overlapping competitors

Player	Overlap	Differentiation
Zapier / n8n / Make	Workflow execution	Tier-1 partner (n8n) + Tier-2 (Zapier); we are upstream discovery layer
Clay / Apollo	GTM data	We are one layer below; Clay uses providers we recommend
Hugging Face	Discovery hub	Different artifact (models vs agent orchestrations); HF could pivot but their DNA is artifact storage
G2 / Capterra	Vendor-neutral marketplace	Different supply ecosystem (SaaS vs agents); they don't reach this category
Cursor / Cline / Continue	Native MCP wiring	Tier-2 partners — they integrate our recipes
LangChain / LlamaIndex / AutoGen	Component discovery	Component-level; we are workflow + recipe + marketplace
Frontier model vendors	Native tool curation	Vendor-neutral, multi-host, multi-protocol; they are Tier-1 partners

10.3 The credible threats (and our defenses)

Threat	Why credible	Defense
Anthropic launches a ranked, benchmarked MCP marketplace	Owens protocol; has distribution	Vendor-neutral, multi-host, multi-protocol; A2A and AI-native they won't include; benchmark methodology open and replicable; recipe handoff to every host. Plus Anthropic is our Tier-1 partner — co-opt the threat.
Hugging Face extends to agents	Brand, community	Agent workflows ≠ model artifacts; recipe handoff is workflow-shaped; trust ladder built for agent risks (auth, side effects, cost)
Smithery raises and adds rankings + marketplace	Already has the index	Multi-source dedupe across 17+ sources spanning canonical registries, third-party aggregators, package registries, and community awesome-* lists; goal decomposition + recipe handoff; vendor-neutrality vs single-registry capture

Threat	Why credible	Defense
Microsoft / GitHub launches an agent registry	Distribution + IDE integration	Cross-host neutrality; benchmark + verification not as easily replicated as a catalog
n8n in-sources discovery	They could, eventually	Tier-1 partnership locks them in; our cross-host coverage (Cursor + Cline + Claude + Zapier + own apps) is broader than n8n's reach

11. Moat and defensibility

11.1 Six compounding assets

Asset	What compounds	Years to replicate
Discovery freshness	17+ source scouts across 4 channel classes, daily refresh	2–3
Verification evidence	Trust tier ladder, MCP probe history, vendor cross-refs	3–5
Benchmark history	Per-capability runs, time-series, methodology trust	3–5
Refusal demand signal	Every unmet goal → demand event → public leaderboard	Not replicable without our distribution
Marketplace network effects	Vendors register because users are here; users come because vendors are verified; both sides reinforce	Compounds with both sides
Tier-1 partnership distribution	Each integration (n8n, Cursor, Anthropic) creates switching cost for the partner	Once locked, hard to replicate

11.2 Network effects (multi-directional)

- More users → more vendor demand to be listed → vendor revenue → more benchmark data → better recipes → more users
- More vendors registered → more credible recipes → more LLM hosts integrate → more users → more vendor registrations
- More benchmark runs → more credible rankings → more vendor-neutral buyer trust → more enterprise pull
- More partnerships → more distribution → more users → more vendor pull → more partnership leverage

11.3 Switching costs (Pro + Vendor + Partner tiers)

- Saved recipes (team library)

- Private notes per agent
- Freshness alerts wired to user's workflow
- Benchmark API keys wired into vendor product backends
- **Vendor claimed profiles** (analytics, community responses, benchmark history)
- **Sponsored placement contracts** (multi-quarter)
- **Partner integrations** (n8n templates referencing our index)
- Enterprise registry (deep integration with customer's CI/CD or LLMOps)

11.4 Why this is venture-defensible

A trusted, ranked, benchmark-rich, marketplace-powered, partnership-distributed discovery layer over a 10x supply ecosystem (vs models, packages, or SaaS) supports valuations in the \$1–5B range based on direct precedents (Hugging Face, G2, Gartner Digital Markets). The moat compounds with index size, marketplace density, partnership lock-in, and benchmark coverage.

12. Product roadmap (24 months)

12.1 Q1 (now → +3 months)

-  Generic protocol adapters (MCP / A2A / OpenAPI) — done
-  Verification tier ladder + qualification gate — done
-  Goal decomposer + per-capability provider partition — done
-  Pre-plan discovery + post-goal refresh — done
-  Recipe export endpoint (`/recipe/export?format=...`)
-  Public launch — HN, dev Twitter, blog
-  Index density push — 2,000+ verified agentic candidates
-  **n8n + Cursor + Anthropic Tier-1 BD outreach started**

12.2 Q2 (+3 → +6 months)

- Pro tier MVP (saved recipes, teams, freshness alerts)
- Auth (Clerk + Stripe billing)
- BYO-credentials sandbox alpha
- First publishable benchmark cell (web_search likely)
- 5 dev team design partners
- 10K weekly active developers
- **n8n Tier-1 partnership signed and integration live**
- **Cursor partnership LOI signed; "Add to Cursor" button on first recipes**
- **Founding marketplace cohort outreach started** (first 3 LOIs target)

12.3 Q3 (+6 → +9 months)

- Benchmark API beta
- 3 publishable benchmark cells
- First enterprise design partner (private registry pilot)
- GTM data category content expansion
- **Cursor Tier-1 partnership signed and integration live**
- **First 3 founding cohort Verified Benchmarks published**

- 25K weekly active users

12.4 Q4 (+9 → +12 months)

- Benchmark API GA
- Enterprise registry private beta
- **Vendor portal alpha** (claimed profiles + first 5 paying vendors)
- **Anthropic methodology partnership announced**
- **Optional execution fee (tier 8) opens to founding cohort first (1-year waiver); first 1–2 vendors opt in**
- Series A readiness gate: \$50K MRR, 25K WAU, 3 enterprise pilots, 3 publishable benchmark cells, **2 of 3 Tier-1 partner integrations live (n8n + Cursor)**, 5 paying marketplace vendors

12.5 Year 2

- Q1: **Vendor portal GA** (25 paying vendors). **Anthropic "Tested by PlanMyAgents" badge live in Claude Desktop.**
- Q2: **Verified Benchmark Badge program launched publicly.** Cline + Continue.dev Tier-2 partnerships signed. **Optional execution fee (tier 8) opens to all vendors (opt-in, default off).**
- Q3: First sponsored category placements (with disclosure). First enterprise registry pilot closed paying. **First execution-fee vendors generating \$5K–25K MRR.**
- Q4: **Demand-Data API beta.** \$5M ARR run rate. 100 paying vendors. **All 3 Tier-1 integrations live (n8n + Cursor + Anthropic).** Series B readiness.

12.6 Year 3+

- Cloud OEM (AWS / Azure / GCP)
- International (EU, APAC) hosting + i18n
- Recipe marketplace (community-contributed verified recipes, free + curated paid)
- Standards body influence (MCP WG, A2A WG implementer-of-record)
- Acquisition optionality (G2, Gartner, GitHub, AWS, Anthropic)

13. Financial plan

13.1 Use of seed funds (\$4M model)

Category	%	\$
Engineering (3 backend + 2 frontend + 1 ML + 1 marketplace eng)	55%	\$2.2M
Discovery + benchmark infrastructure	15%	\$600K
GTM (DevRel + partnership BD + content + design partners)	20%	\$800K
Ops (legal, finance, runway buffer, SOC2 prep)	10%	\$400K

13.2 Headcount plan

Role	Q1	Q2	Q3	Q4	Y2
------	----	----	----	----	----

Role	Q1	Q2	Q3	Q4	Y2
Founder / CEO	1	1	1	1	1
Backend engineer	0	1	2	3	4
Frontend engineer	0	1	1	2	3
ML engineer	0	0	1	1	1
Marketplace engineer	0	0	0	1	1
DevRel	0	0	1	1	1
Partnership BD	0	1	1	1	2
Sales (enterprise)	0	0	0	0	1
Ops / finance	0.25	0.25	0.5	0.5	1
Total	1.25	4.25	7.5	10.5	15

13.3 Cost structure assumptions

- LLM API spend (Groq, OpenAI, Anthropic): ~\$3K / mo at 10K WAU; ~\$15K / mo at 100K WAU
- Embedding compute: ~\$500–2K / mo
- Postgres + pgvector: ~\$300–1K / mo
- Hosting: ~\$500–2K / mo
- **Benchmark runs (recurring + vendor-paid): ~\$5K / mo at scale**
- Total infra at Year 2 with 50K WAU + 100 paying vendors: ~\$30K / mo → \$360K / yr

13.4 Revenue projection (conservative)

Quarter	WAU	Pro MRR	Bench API MRR	Marketplace (4-6) MRR	Enterprise MRR	Exec Fee (tier 8) MRR	Total MRR
Q1	1K	\$0	\$0	\$0	\$0	\$0	\$0
Q2	5K	\$3K	\$0	\$0	\$0	\$0	\$3K
Q3	15K	\$12K	\$5K	\$0	\$0	\$0	\$17K
Q4	25K	\$25K	\$15K	\$5K (5 vendors)	\$10K (1 pilot)	\$0 (waiver yr)	\$55K
Y2 Q1	35K	\$40K	\$30K	\$15K (25 vendors)	\$30K	\$0 (waiver yr)	\$115K
Y2 Q2	45K	\$55K	\$45K	\$35K (50 vendors)	\$60K	\$5K (5 opt-ins)	\$200K
Y2 Q3	55K	\$70K	\$60K	\$60K (75 vendors)	\$120K	\$15K (8 opt-ins)	\$325K

Quarter	WAU	Pro MRR	Bench API MRR	Marketplace (4-6) MRR	Enterprise MRR	Exec Fee (tier 8) MRR	Total MRR
Y2 Q4	65K	\$90K	\$80K	\$90K (100 vendors)	\$250K	\$30K (10 opt-ins)	\$540K
Y3 Q4	100K	\$150K	\$200K	\$300K (300 vendors)	\$1.2M	\$180K (60 opt-ins)	\$2.03M

→ Year 3 ARR ~ **\$24M** on conservative assumptions (the tier-8 execution fee adds ~10% on top by Y3 end).

13.5 Series A target (12–15 months out)

- \$50–100K MRR
- 25K+ weekly active developers
- 3+ enterprise pilots (1 closed paying)
- 3+ publishable benchmark cells
- **2 of 3 Tier-1 partner integrations live (n8n + Cursor); Anthropic methodology announced**
- **5+ paying marketplace vendors; 1–2 founding cohort vendors opted into execution fee (waiver year)**
- Brand recognition in dev community

Target raise: \$10–20M Series A at \$50–100M post.

14. Team and hiring

14.1 Founder

Deepraj Jha — CEO / founding engineer.

- 11 years backend engineering (Adobe, Amazon Payments, Salesforce)
- Production systems at scale; payment + cost discipline; external API integration depth
- Vendor-neutral by background

14.2 Founding hires (in order)

1. **Senior backend engineer** — discovery infra, scout fleet, store performance
2. **Partnership BD lead** — Tier-1 partnership deals (n8n, Anthropic, founding vendors)
3. **Senior full-stack engineer** — Pro tier UX, recipe export, dashboard polish
4. **ML engineer** — embedding pipeline, benchmark methodology, judge model tuning
5. **DevRel / community lead** — public content, dev community, design partner cohort
6. **Marketplace engineer** — vendor portal, claimed profile, sponsorship surface
7. **Sales (first enterprise)** — only after 3 inbound enterprise pilots prove pull

14.3 Advisor profile we are recruiting

- One MCP/A2A protocol insider (Anthropic, Google, ServiceNow, Salesforce)
- One marketplace operator (G2, Capterra, Hugging Face, Replicate, Snyk)
- One distribution / DevRel veteran (Vercel, Supabase, Replicate, Hugging Face, Hashicorp)

- One enterprise GTM operator (Gartner, G2, Snyk, Datadog)
- One ML benchmark / evaluation researcher (academic credibility)

15. Risks and mitigations

Risk	Likelihood	Impact	Mitigation
Frontier model vendor launches competing registry	Medium	High	Vendor-neutrality, multi-protocol, multi-host; co-opt via Tier-1 Anthropic partnership
MCP/A2A adoption stalls	Low	Existential	Multi-protocol coverage including AI-native services
Discovery quality degrades	Medium	High	Continuous freshness monitoring, multi-source dedupe, verification ladder, judge oversight
LLM cost explodes	Medium	Medium	Escalating client with tiered fallback; embedding caching; judge model rotation
Slow Pro tier conversion	Medium	Medium	Free tier itself is acquisition motor; Pro features are usage-driven
Slow vendor-side conversion	Medium	Medium	Founding vendor cohort with discounted founding pricing; case studies
Enterprise sales cycle longer than expected	High	Medium	Design partner program de-risks before scaling sales
Vendor pushback on benchmark publication	Medium	Low	Transparent methodology, right-of-reply, vendors can re-run
Vendor-neutrality firewall breach	Low	Existential	Codified in benchmark methodology + pricing; would require explicit policy violation
Tier-1 partner deals fail to materialize	Medium	Medium	Multiple parallel BD tracks; n8n + Anthropic are the priority but not the only candidates
Compliance / privacy issue	Low	High	We don't store customer credentials; minimal PII; SOC2 path from Q3
Founder burnout / single point of failure	Medium	High	Co-founder search post-seed; advisor bench
Series A fails despite hitting metrics	Low	High	Default-alive on Pro + Benchmark API + early Marketplace by month 18 with conservative spend

16. Fundraising narrative

16.1 The one-sentence pitch

"PlanMyAgents is the trust, recipe, and marketplace platform for the AI agent economy — G2 + Hugging Face + Gartner Digital Markets, for an ecosystem 10x the surface area of any of them."

16.2 Why now (one paragraph)

MCP shipped 14 months ago; A2A shipped 9 months ago; every LLM client now has native tool use; n8n / Zapier / Make are racing to add AI nodes they don't have a discovery layer for. The supply of agents is exploding past 1,000 MCP servers, hundreds of A2A agents, and thousands of AI-native services — but nothing ranks them, verifies them, benchmarks them, or runs a vendor-neutral marketplace on top. Every dense supply ecosystem (packages, models, SaaS, apps) produced a discovery + marketplace layer worth billions; this ecosystem will too. We are building it before a frontier model vendor or marketplace incumbent realizes the layer is missing.

16.3 Why us (one paragraph)

11 years backend engineering at Adobe, Amazon Payments, and Salesforce — the right discipline mix for an index that has to be fresh, trustworthy, cost-disciplined, and enterprise-credible. Vendor-neutral by background. Already shipped a working prototype with 14+ source connectors, semantic decomposition, verification tier ladder, generic protocol adapters, and 1,200+ backend tests passing.

16.4 Why this round (one paragraph)

\$3–5M seed funds 18 months of: shipping the recipe export and reaching 25K weekly active developers, signing **Tier-1 partnerships with n8n + Cursor + Anthropic** (workflow distribution + IDE / coding host distribution + protocol authority), running a parallel **founding marketplace cohort workstream** to land 5–10 paying vendors and the first 3 publishable Verified Benchmarks, **piloting the optional vendor-opt-in execution fee (tier 8)** with the founding cohort under a 1-year waiver, closing 3 enterprise design partners, and reaching \$50K MRR — the Series A gate. The seed is a capital-efficient launch, not a multi-year platform build; the marketplace + partnership flywheels compound from Year 1 end onward.

16.5 Why this investor profile

We are looking for one or two leads with:

- Infrastructure / developer tools investing track record (Vercel, Supabase, Hugging Face, Replicate, Hashicorp, Datadog, etc.)
- **Marketplace investing experience (G2, Capterra, Etsy, Faire, etc.)**
- Comfort with usage-driven seed motion + marketplace flywheel thesis
- Belief that benchmark + trust + neutrality can be a moat (not connector count)
- Willingness to introduce us to **n8n, Cursor, Anthropic, Cline, Continue.dev, Google (A2A), AWS Bedrock, GCP Vertex, and Tier-1 vendor agent companies (Apollo, Perplexity, Hunter, Linkup, Tavily, Exa, Apify)** for ecosystem partnerships

17. Appendix: glossary and FAQ

Glossary

- **MCP (Model Context Protocol)** — open standard from Anthropic for LLM tool integration; primary agent supply we index
- **A2A (Agent-to-Agent)** — open protocol for agent interop; secondary supply
- **AI-native service** — API service with agent-shaped interfaces (Perplexity, Linkup, Exa, Apify, etc.); tertiary supply
- **Capability** — atomic, namable thing an agent can do ([web_search](#), [payment_authorization](#), etc.); 26 in routable catalog plus open-ended labels
- **Trust ladder** — five-tier evidence model: [capability_verified](#) > [registered_in_directory](#) > [known_provider](#) > [community_listed](#) > [unverified](#)
- **Recipe** — machine-and-human-readable workflow artifact the user executes in their own LLM host
- **Goal decomposer** — LLM that turns user goal into ordered sub-tasks with capability slugs
- **Discovery gap event** — recorded signal whenever a goal sub-task has no qualified agent; fuels public gap leaderboard
- **BYO-credentials sandbox** — optional execution surface where user pastes credentials in browser session; we invoke MCP/A2A/OpenAPI on their behalf; credentials never persisted server-side
- **Vendor portal (Phase 2)** — vendor self-serve surface for claimed profile, verified benchmark, sponsored placement, lead routing, demand-data API
- **Verified Benchmark Badge (Phase 2)** — extended benchmark run (≥ 250 samples) result published whatever it is; vendor pays for the run, not the result
- **Vendor-neutrality firewall** — money buys *visibility* (sponsored placement, claimed profile), never *ranking position*; codified in pricing + methodology
- **Tier-1 partnership (Phase 3)** — founding-priority partner (n8n, Anthropic, founding vendors) where the cost of inaction is highest

FAQ

Q: Why not just be Zapier for AI agents? A: Zapier has 5,000+ hand-coded connectors and a \$5B valuation. The arms race is over. Trust + benchmark + recipe + marketplace is unfilled and has \$25B+ comparable valuations (Gartner). And Zapier is our **Tier-2 partner**, not our competitor.

Q: Why won't Anthropic or OpenAI just do this? A: They might curate within their own LLM host. Vendor-neutral, multi-protocol, multi-host benchmarking with recipe handoff across every host is a different product. Their incentive is single-host capture; ours is to be the meta-layer. **Anthropic is our Tier-1 partner.** Co-opt the threat.

Q: Won't Hugging Face extend to agents? A: They could. HF Spaces is the closest extension and has been stalled. HF's DNA is artifact storage; agent workflows are a different shape, with different benchmarking methodology, different trust signals, and a recipe handoff that crosses LLM hosts. If they extend, they'll do MCP only — we still own A2A + AI-native + cross-host.

Q: Won't Smithery just add ranking and become a marketplace? A: They could try. Multi-source dedupe across 14+ sources is hard to replicate (they're 1 source). Goal decomposition + recipe handoff is meaningfully proprietary. Vendor-neutrality is a *positioning* choice that, once compromised, can't be uncompromised. Our window to be *the* trusted marketplace is now.

Q: How do you measure "best agent" objectively? A: Documented benchmark methodology per capability ([docs/benchmark-methodology.md](#)): success rate against known good answers, p95 latency,

cost per query, recall@K for search, capability-specific quality measures. Methodology is public so vendors can dispute and re-submit.

Q: What if vendors won't be benchmarked? A: We benchmark from public endpoints with documented inputs; vendors are notified before publication; we offer right-of-reply and re-run on request. Vendors who refuse get "no benchmark" rather than a fabricated score. Our **Verified Benchmark Badge** is opt-in *and* the result is published whatever it is — no hiding poor results.

Q: How does the marketplace stay vendor-neutral if vendors pay you? A: Money buys *visibility* (sponsored placement, claimed profile features, lead routing) and optionally *execution attribution* (tier 8). Money never buys *ranking position*. This firewall is codified in pricing + methodology + product surface. Tier 8 (the opt-in execution fee) is hard-capped at 2% and architecturally decoupled from the ranker — vendor cannot pay a higher rate to influence position. G2 has held this line for 14 years; it's the asset.

Q: How is the optional execution fee different from a Zapier-style take rate? A: Five concrete differences. (1) **Vendor pays, not user** — vendor settles out of their margin; user price never changes. (2) **Opt-in per vendor, default off** — no vendor pays unless they signed the addendum. (3) **Hard-capped at 2%** by contract — vendor cannot escalate. (4) **Decoupled from ranker by architecture** — the ranking module does not import the marketplace store; CI lints enforce this. (5) **Only applies when our BYO-creds sandbox is in the runtime path** — a user running the same recipe in Claude Desktop, Cursor, or n8n directly triggers no fee. Comparable is Amazon Associates attribution, not Stripe / Zapier / App Store.

Q: How big can this really get? A: Hugging Face is \$4.5B post on models alone. G2 is ~\$2B+ on SaaS. Gartner is \$25B mcap on enterprise IT advisory. The agent supply ecosystem is 10× more granular and fast-moving. A trusted discovery + marketplace + partnership platform here is plausibly a \$5–20B outcome.

Q: Why open this now vs after a closed launch? A: The supply is mid-explosion. Every quarter delays the index density flywheel and the marketplace network effect. The first credible discovery + marketplace layer captures category-defining brand value.

Q: What's the acquisition path? A: G2 (extending into agents), Gartner Digital Markets (extending category coverage), GitHub / Microsoft (extending Copilot Workspace), AWS (extending Bedrock developer experience), Anthropic (folding into MCP), Google (folding into A2A). We optimize for product-independence; acquisition optionality is a byproduct.